

HW 4B: Stochastic Gradient Descent and Lipschitz Extensions

CS 208 Applied Privacy for Data Science, Spring 2019

Version 1.0: Due Tuesday, April 30, 11:59pm.

Instructions: Submit a single PDF file containing your solutions, plots, and analyses. Make sure to thoroughly explain your process and results for each problem. Also include your documented code and a link to a public repository with your code (such as GitHub/GitLab). Make sure to list all collaborators and references.

- For each of the following sets \mathcal{G} of datasets and neighbor relations \sim , hypotheses $\mathcal{H} \subseteq \mathcal{G}$, and functions $f : \mathcal{G} \rightarrow \mathbb{R}$, calculate (i) the global sensitivity of f (denoted GS_f or ∂f), (ii) the minimum local sensitivity of f , i.e. $\min_{x \in \mathcal{G}} \text{LS}_f(x)$, and (iii) the restricted sensitivity of f (denoted $\partial_{\mathcal{H}} f$ or $\text{RS}_f^{\mathcal{H}}$). For Part 1a, also describe an explicit Lipschitz extension of f from \mathcal{H} to all of \mathcal{G} .
 - $\mathcal{G} = \mathbb{R}^n$ where $x \sim x'$ if x and x' differ on one row, $\mathcal{H} = [a, b]^n$ for real numbers $a \leq b$, and $f(x) = (1/n) \sum_{i=1}^n x_i$.
 - $\mathcal{G} = \mathbb{R}^n$ where $x \sim x'$ if x and x' differ on one row, $\mathcal{H} = [a, b]^n$ for real numbers $a \leq b$, and $f(x) = \text{median}(x_1, \dots, x_n)$.
 - \mathcal{G} = the set of undirected graphs (without self-loops) on vertex set $\{1, \dots, n\}$ where $x \sim x'$ if x and x' are identical except for the neighborhood of a single vertex (i.e. node privacy), \mathcal{H} the set of graphs in \mathcal{G} in which every vertex has degree at most d for a parameter $2 \leq d \leq n-1$, and $f(x)$ = the number of isolated (i.e. degree 0) vertices in x .
- Recall for data universe $\mathcal{X} = \{x \in \mathbb{R}^k : \|x\| \leq R\}$ and $\varepsilon < 1$, the following is a (ε, δ) -DP local randomizer:

$$Q_{\text{gauss}}(x) = x + \mathcal{N}\left(0, \left(\frac{R}{2\varepsilon^2}\right)^2 \cdot \ln\left(\frac{1.25}{\delta}\right) \cdot I_k\right).$$

We mentioned in class that there is also a pure DP randomizer that can be used instead. It works as follows.

$Q_{\text{pure}}(x)$ for $x \in \mathcal{X}$:

- Choose a uniformly random unit vector $u \in \{v \in \mathbb{R}^k : \|v\| = 1\}$.
- Do randomized response on $\langle u, x \rangle / R \in [-1, 1]$ to obtain a binary value $b \in \{\pm 1\}$.
- Output bu

Analyze Q_{pure} as follows.

- Prove that Q_{pure} is ε -DP.
- Argue that for every $x \in \mathcal{X}$, we have

$$\mathbb{E}[Q_{\text{pure}}(x)] = \frac{e^\varepsilon - 1}{e^\varepsilon + 1} \cdot \frac{c_k}{R} \cdot x,$$

for a constant c_k that depends only on the dimension k . (Hint: decompose $u = u^\parallel + u^\perp$ where u^\parallel is a scalar multiple of x and u^\perp is orthogonal to x .) You don't need to calculate the constant c_k (but it happens to equal $2s_{k-1}/((k-1) \cdot s_k)$, where s_d denotes the $(d-1)$ -dimensional volume of the unit sphere $\{v \in \mathbb{R}^d : \|v\| = 1\}$ in d dimensions).

(c) Thus, like $Q_{\text{gauss}}(x)$, the post-processed output

$$Q'_{\text{pure}}(x) = \frac{e^\varepsilon + 1}{e^\varepsilon - 1} \cdot \frac{R}{c_k} \cdot Q_{\text{pure}}(x)$$

is an unbiased estimator of x . Ignoring constant factors, assuming $\varepsilon \leq 1$, and using the fact that $c_k = \Theta(1/\sqrt{k})$, compare the maximum expected squared error $\max_{x \in \mathcal{X}} \mathbb{E}[\|Q(x) - x\|^2]$ for $Q = Q'_{\text{pure}}$ vs. $Q = Q_{\text{gauss}}$.

3. In our code example,¹ we saw how to release an estimated Logistic regression using differentially private stochastic gradient descent (DP-SGD) to optimize a regression under the centralized model. Convert this code to once again release the probability of marriage given education level, but using DP-SGD under the *local* model. Recall that local DP does not satisfy privacy amplification by subsampling, but you can achieve a similar effect by rotating through disjoint batches, so that each individual participates in at most $\lceil T \cdot B/n \rceil$ batches, where T is the number of iterations and B is the batch size.² Using Gaussian noise as in the current code will give an (ε, δ) local DP implementation. Also construct a pure $(\varepsilon, 0)$ local DP implementation using Problem 2, and compare the performance of the two implementations.

¹See https://github.com/privacytoolsproject/cs208/blob/master/examples/wk7_localmodel/privateSGD.r and `privateSGD.ipynb`.

²Note, in the code example, $T = 1$, $B = \sqrt{n}$.